

Business Recommendations

A concise decision brief for HR and recruiting leaders. The system is designed to prioritise recruiter attention, not to replace human hiring judgement.

RECOMMENDED DECISION

Use the tool to sort HR candidates for recruiter review: begin with the strongest text matches, then let recruiter feedback refine the order.

Recommendation at a glance

1

Remove duplicates before ranking

Rank each unique candidate once so repeated source rows do not distort the shortlist.

2

Start with a consistent initial ranking

Use candidate title and career-intent text to create the first recruiter review order.

3

Let recruiter feedback refine the order

A star or reject can move similar profiles up or down for the current search.

4

Keep the recruiter in control

Use the ordering to focus attention; do not use it to make shortlist or hiring decisions automatically.

Evidence base

104

SOURCE ROWS

The source file includes repeated records.

53

UNIQUE PROFILES

Independent evaluation is based on unique candidate profiles.

192

SETUPS COMPARED

Many reasonable ranking approaches were compared systematically.

0.960

TOP-10 QUALITY

The selected initial ranking scored very strongly in grouped validation.

What this means for the business

- The system should prioritise recruiter attention; it should not make hiring decisions.
- The evidence supports a simple, text-driven starting order rather than extra scoring complexity.
- Better recruiter feedback and hiring-outcome data are the most valuable next steps.

How the ranking works

There is one human-labelling step and two ranking stages. The same 0-3 relevance labels are used to build and test the initial order; recruiter feedback is then used only to refine that order for the current search.

How candidate relevance was labelled

Each of the 53 unique profiles received one human relevance grade from 0 to 3. The rules considered both HR relevance and whether the profile reflected the approved search intent (for example, aspiring or seeking HR). Keyword presence alone was not enough, and duplicate source rows were not treated as independent judgements.

0**Irrelevant**

No meaningful HR relevance to the search objective.

1**Weak / adjacent**

Some HR or career overlap, but materially mismatched intent.

2**Clearly HR-relevant**

Relevant HR profile, but not the strongest direct target.

3**Strong direct match**

Clear HR relevance plus strong aspiring/seeking intent.

The two ranking stages

STAGE 1 - INITIAL RANKING**PURPOSE**

Create a consistent first-pass order before the recruiter gives any feedback.

MODEL CHOSEN

A simple regularised linear scorer (Ridge) using several text-similarity signals: word/character TF-IDF, BM25 and direct phrase overlap. Connection count was not retained.

WHY THIS MODEL

It performed strongly while remaining simple, stable and easier to explain.

STAGE 2 - FEEDBACK REFINEMENT**PURPOSE**

Personalise the current search after the recruiter stars or rejects profiles.

MODEL CHOSEN

A classic relevance-feedback method (Rocchio) that moves the search toward language found in preferred profiles and away from language found in rejected profiles.

HOW IT IS CONTROLLED

The refined order keeps 65% of the original ranking and uses 35% feedback, so feedback adjusts rather than replaces the validated baseline.

How we checked ranking quality

0.960**NDCG@10**

Did the strongest human-rated candidates appear near the top? Near 1.0 means the top of the list closely matched the human relevance grades.

1.000**MAP@10**

Did relevant candidates consistently appear within the first ten positions? 1.0 means they did so perfectly in validation.

1.000**MRR**

How quickly did the first relevant candidate appear? 1.0 means the first position was relevant in validation.

These are offline ranking-quality measures from duplicate-safe grouped validation; they are not measures of hiring success or employee performance.

What the organisation should do

Four practical actions, ordered by importance. The emphasis is a controlled pilot with clear decision boundaries, not automated hiring.

Priority actions

1

Deduplicate before ranking

WHY IT MATTERS

Repeated source rows can make the evidence look stronger than it is.

BUSINESS OUTCOME

Cleaner shortlists and more defensible evaluation.

2

Use the validated initial ranking

WHY IT MATTERS

The selected text-based model gives recruiters a consistent first-pass review order.

BUSINESS OUTCOME

A simple, explainable and easier-to-govern starting point.

3

Use recruiter feedback to refine the list

WHY IT MATTERS

Stars and rejects can adjust the current search without replacing the original ranking.

BUSINESS OUTCOME

A more responsive shortlist while preserving a stable reference.

4

Collect better recruiter and outcome data

WHY IT MATTERS

The current evidence base is 53 uniquely assessed candidate profiles.

BUSINESS OUTCOME

Stronger future testing and safer expansion to new searches.

Decision boundaries

USE THE SYSTEM FOR

- Prioritising which HR profiles a recruiter reviews first.
- Refining a search after explicit recruiter feedback.
- Providing a consistent first-pass candidate order across repeated reviews.

DO NOT USE THE SYSTEM FOR

- Automatically deciding who should be hired.
- Treating the ordering as a prediction of job success.
- Using connection count as a hiring-quality proxy.

Pilot, measure, then scale

A controlled rollout keeps the system measurable and under recruiter oversight while the evidence base grows.

Recommended rollout

1

Pilot

2-4 WEEKS

- Use on the current HR search only.
- Keep recruiter review as the final decision point.
- Capture star/reject feedback and shortlist changes.

2

Measure

4-8 WEEKS

- Compare the ordered candidate lists with recruiter judgement.
- Track review time and shortlist relevance.
- Review false positives and missed strong candidates.

3

Scale

ONGOING

- Expand only after stable pilot evidence.
- Refresh relevance judgements when role intent changes.
- Monitor shortlist quality and recruiter adoption.

Success measures

- Recruiter review time per search decreases.
- Top-ranked candidates are judged relevant more often.
- Fewer duplicate profiles distort the shortlist.
- Recruiters use feedback controls consistently.
- No material decline in quality when the search changes.

Governance controls

- Human recruiter remains accountable for shortlist and hiring decisions.
- Keep the original ranking available when feedback adjusts the list.
- Exclude connection count from the selected ranking approach.
- Log feedback and ranking changes for review.
- Reassess labels and performance before expanding to materially different roles.

BOTTOM LINE

Deploy the ranking as a recruiter decision-support tool: use the validated initial ordering to focus attention, use recruiter feedback only to refine the current search, and expand only after collecting stronger real-world outcome evidence.